

Lab 03

ENVX2001 Applied Statistical Methods

Contents

Welcome!	2
Learning outcomes	2
Specific goals	2
Preparation	2
Downloads	3
1. Exploring diatom diversity (~15 min)	3
Getting started	3
Worked Example 1	3
Exploring the data	4
Exercise 1	4
.....	5
Exercise 2	5
2. Testing for differences (~25 min)	5
The ANOVA model	5
Fitting the model	5
Worked Example 2	6
Interpreting the results	6
QQ plot	6
Shapiro-Wilk test	7
Exercise 3	7
Post-hoc comparisons	8
Worked Example 3	8
.....	10
Exercise 4	10
3. Comparing <i>t</i> -tests and ANOVA (~20 min)	10
The data	10
.....	11
Exercise 5	11
.....	11
Exercise 6	11
4. Lambs (~25 min)	11
The data	11
.....	12
Exercise 7	12
Conclusion	12
Closing thoughts	12

Welcome!

This week we will analyse diatom diversity in streams, compare comb weights in chicks, and test immunoglobulin levels across lamb breeds using one-way ANOVA.

You will come across two types of tasks in this lab. **Worked Examples** have solutions you can expand and check straight away. **Exercises** do not – solutions for these will be posted on Friday evening.

As always, don't forget to take regular breaks in between sections. A good rule of thumb is to work for 20-25 minutes, then take a 5-minute break. This will help you stay focused and avoid burnout.

Learning outcomes

In this lab, we will learn how to:

1. Explore grouped data using summary statistics and graphical summaries.
2. Fit a one-way ANOVA using `aov()` and interpret the ANOVA table.
3. Perform post-hoc comparisons using `emmeans`.
4. Explain when a two-sample *t*-test gives the same results as a one-way ANOVA.

Specific goals

By the end of this lab, you should be able to:

- ☐ Convert character variables to factors using `as.factor()`
- ☐ Calculate group means and standard deviations
- ☐ Produce boxplots for grouped data
- ☐ Fit a one-way ANOVA and read the output table
- ☐ Use `emmeans` for post-hoc testing
- ☐ Compare *t*-test and ANOVA results for a two-group design

Preparation

This lab uses `tidyverse`, `readxl`, and `emmeans`. Install any you are missing by running the following **in the console**:

CODE

```
install.packages(c("tidyverse", "readxl", "emmeans"))
```

Downloads

File	Used in	Download
diatoms.xlsx	Sections 1–2	Download
chick_marigold.xlsx	Section 3	Download
lambs.csv	Section 4	Download

Save all files into a folder called data inside your project folder.

1. Exploring diatom diversity (~15 min)

Coal miners once carried canaries underground as indicators of air toxicity. Diatoms are the freshwater equivalent, acting as bioindicators whose diversity changes with water quality. Medley and Clements (1998) surveyed 34 streams in the Rocky Mountains of Colorado, where decades of mining had left zinc leaching into the water, and counted diatom species at each site. Does zinc contamination reduce diatom diversity?

Getting started

Let us warm up. This section is a quick refresher on data import, summary statistics, and plotting and should be straightforward if you have been following the labs in the past couple of weeks. If you have any questions, ask your tutor or post on Ed Discussion.

Worked Example 1

Import the Diatoms worksheet from `diatoms.xlsx`, check the data structure with `str()`, and convert the character variables `Zinc` and `Stream` to factors. Use the code chunk below as a template for your own report.

If you have difficulty importing, see the [data import guide](#).

CODE

```
library(tidyverse)
library(readxl)

# Look at the Excel file to work out the worksheet name and range:
diatoms <- read_excel(
  path = "data/diatoms.xlsx",
```

```

    sheet = ...,
    range = ...
  )

  str(diatoms)

# Convert character variables to factors:
diatoms <- diatoms |>
  mutate(
    Zinc = as.factor(Zinc),
    Stream = as.factor(Stream)
  )

```

Solution

```

CODE
library(tidyverse)
library(readxl)

diatoms <- read_excel(
  path = "data/diatoms.xlsx",
  sheet = "Diatoms",
  range = "A1:D35"
)

str(diatoms)

diatoms <- diatoms |>
  mutate(
    Zinc = as.factor(Zinc),
    Stream = as.factor(Stream)
  )

```

Exploring the data

Before running any formal test, it helps to look at summary statistics and plots to see what patterns emerge.

Exercise 1

Calculate the mean and standard deviation of Diversity for each level of Zinc. What do the values suggest about differences between groups?

Even before doing any formal statistics, summary statistics are useful decision tools. If the means are very different and the standard deviations are small, we might already have a good idea that there are differences between groups, and knowing the distribution of the data can help us decide which test to use and how to check assumptions.

However, numbers alone can be hard to interpret. This is where plotting comes in to present the data in a more intuitive way. A boxplot is a good choice for grouped data, as it shows the median, interquartile range, and potential outliers.

Exercise 2

Produce a boxplot of `Diversity` grouped by `Zinc`. Add appropriate axis labels. What does the plot tell you about differences between groups?

With fewer than 10 observations per group, boxplots are often more informative than histograms.

Take this opportunity to explore the data further if you like. You could make a histogram of the `Diversity` values within each `Zinc` group, or calculate other summary statistics such as the range or interquartile range. The more you understand the data before fitting a model, the better you will be able to interpret the results and check assumptions.

Once you are done, take a quick break. It is important to step away from the screen regularly.

2. Testing for differences (~25 min)

We noticed (hopefully) in Section 1 that the `HIGH` zinc group has lower diatom diversity than the others. But is that difference real, or could it be random variation? A one-way ANOVA will test this formally.

The ANOVA model

A one-way ANOVA tests whether the means of several groups are all equal. The model is:

$$y_{i,j} = \mu_i + \varepsilon_{i,j}$$

where $y_{i,j}$ is the diversity for observation j in zinc group i , μ_i is the mean of group i , and $\varepsilon_{i,j}$ is the residual.

The hypotheses are:

$$H_0 : \mu_1 = \mu_2 = \dots = \mu_t$$

$$H_1 : \text{not all } \mu_i \text{ are equal}$$

We will cover model assumptions formally next week when we look at residual diagnostics.

Fitting the model

In general, we “fit a model” when we test hypotheses, because models explain the relationships between data in a manner that is more flexible than fitting mathematical equations.

Worked Example 2

Fit a one-way ANOVA to test whether diatom diversity differs across zinc levels. Use `aov()` to fit the model and `summary()` to view the ANOVA table. Try to interpret the ANOVA table before looking at the solution below.

The general formula is `aov(response ~ group, data = ...)`.

Solution

CODE

```
anova.diatoms <- aov(Diversity ~ Zinc, data = diatoms)
summary(anova.diatoms)
```

OUTPUT

```
          Df Sum Sq Mean Sq F value Pr(>F)
Zinc         3  2.567   0.8555   3.939 0.0176 *
Residuals    30  6.516   0.2172
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The ANOVA table shows:

- **Df**: degrees of freedom. Treatment df = number of groups minus 1 ($4 - 1 = 3$). Residual df = total observations minus number of groups ($34 - 4 = 30$).
- **Sum Sq**: how much variation is explained by the grouping (Zinc) versus left unexplained (Residuals).
- **Mean Sq**: Sum Sq divided by Df. This gives the average variation per degree of freedom.
- **F value**: the ratio of treatment Mean Sq to residual Mean Sq. A large F suggests that group differences are bigger than random noise.
- **Pr(>F)**: the P-value. If this is below 0.05, we reject H_0 .

Interpreting the results

The assumptions of the ANOVA (normality, equal variances, independence) appear reasonable based on our exploration in Section 1. The standard deviations were similar across groups, and nothing in the boxplot suggested serious departures from normality.

However, we should always formally check the assumptions by looking at model residuals. We will cover this in more detail next week, but here is a quick look at normality and equal variance using the approach from Tutorial 3.

QQ plot

CODE

```
qqnorm(residuals(anova.diatoms), main = "")  
qqline(residuals(anova.diatoms), col = "red")
```




Figure 1: Normal QQ plot of residuals from the diatom ANOVA model.

Shapiro-Wilk test

```
CODE  
shapiro.test(residuals(anova.diatoms))
```

```
OUTPUT  
  
Shapiro-Wilk normality test  
  
data:  residuals(anova.diatoms)  
W = 0.96883, p-value = 0.43
```

The QQ plot shows points falling roughly along the line, and the Shapiro-Wilk test is not significant ($P > 0.05$), so normality is reasonable. We already checked the SD ratio in Section 1 and found it was below 2, supporting equal variances.

Exercise 3

Based on the ANOVA output from Worked Example 2, is there a significant difference in diatom diversity across zinc levels? Write a one-sentence reporting statement that includes the F statistic, degrees of freedom, and P-value.

A reporting statement typically follows this pattern: “There was a significant difference in [response] across [groups] ($F = \dots$, $df = \dots$, $P = \dots$).”

Post-hoc comparisons

The ANOVA tells us that at least one group mean differs from the others, but not which ones. We use post-hoc testing to find out.

Worked Example 3

Use the `emmeans` package to compute estimated marginal means and their confidence intervals for each zinc level. Then visualise the results with `plot()`.

`emmeans(model, "factor_name")` takes a fitted model and the name of the grouping variable.

Solution

CODE

```
library(emmeans)
```

OUTPUT

Welcome to emmeans.
Caution: You lose important information if you filter this package's results.
See '? untidy'

CODE

```
posthoc <- emmeans(anova.diatoms, "Zinc")  
posthoc
```

OUTPUT


Zinc	emmean	SE	df	lower.CL	upper.CL
BACK	1.80	0.165	30	1.461	2.13
HIGH	1.28	0.155	30	0.961	1.60
LOW	2.03	0.165	30	1.696	2.37
MED	1.72	0.155	30	1.401	2.04

Confidence level used: 0.95

The output shows the estimated mean diversity for each zinc level, along with the standard error and 95% confidence interval. Groups whose confidence intervals do not overlap are likely to be significantly different.

CODE

```
plot(posthoc)
```



Exercise 4

Based on the emmeans output from Worked Example 3, which zinc levels differ from each other?
How can you tell from the confidence intervals?

Now is a good time to take a 5-minute break.

3. Comparing t -tests and ANOVA (~20 min)

So far we have used ANOVA on data with four groups. But what happens when there are only two? In Lecture 03a we learned the two-sample t -test, and in Lecture 03b we learned ANOVA. This section shows they give the same answer when there are two groups.

An experiment compared 15-day mean comb weights (g) of male chicks receiving one of two hormones: A (testosterone) or C (dehydroandrosterone). With only two groups, we could use either a pooled two-sample t -test or a one-way ANOVA.

The data

Read the *Comb* worksheet from *chick_marigold.xlsx* and convert *Hormone* to a factor:

CODE

```
comb <- read_excel("data/chick_marigold.xlsx", sheet = "Comb") |>
  mutate(Hormone = as.factor(Hormone))
comb
```

OUTPUT

```
# A tibble: 22 × 2
  CombWt Hormone
  <dbl> <fct>
1    57 A
2   120 A
3   101 A
4   137 A
5   119 A
6   117 A
7   104 A
8    73 A
9    53 A
10   68 A
# i 12 more rows
```

Exercise 5

Before running any test, check whether the assumptions of normality and equal variance are reasonable. Produce a boxplot and calculate the standard deviations for each hormone group.

Exercise 6

Run a pooled two-sample t -test using `t.test(..., var.equal = TRUE)` and a one-way ANOVA using `aov()` on the same data. Compare:

- i. the degrees of freedom,
- ii. the P-values, and
- iii. the relationship between the t statistic and the F statistic.

When there are two groups, $F = t^2$. Verify this for your results.

Use `var.equal = TRUE` in `t.test()` to get the pooled (Student's) t -test, which matches the ANOVA assumption of equal variances.

Take a break!

4. Lambs (~25 min)

This final section is a capstone exercise. You will work through the entire ANOVA workflow on a new dataset with minimal guidance, bringing together everything from Sections 1 to 3.

The levels of immunoglobulin (Ig) in blood serum (g/100 ml) in three breeds of newborn lambs have been investigated. A total of 44 lambs were sampled, with approximately equal numbers per breed. Does immunoglobulin level differ between breeds?

The data

Here is a preview of the dataset. You will import it yourself as part of the exercise.

```
OUTPUT
# A tibble: 44 × 2
  Ig Breed
<dbl> <fct>
1  1.1 1
2  2.2 1
3  1.7 1
```

```
4  1.4 1
5  1.6 1
6  2.3 1
7  1.4 1
8  1.9 1
9  0.8 1
10 1.6 1
# i 34 more rows
```

Exercise 7

Analyse the `lambs` dataset using the full ANOVA workflow. Test the hypothesis that:

$$H_0 : \mu_1 = \mu_2 = \mu_3$$

$$H_1 : \text{not all } \mu_j \text{ are equal}$$

where μ_j is the mean Ig level for breed j .

Work through each step:

- ☐ Read the data and convert `Breed` to a factor.
- ☐ Produce a boxplot and calculate group means and standard deviations.
- ☐ Check the equal variance assumption using the SD ratio.
- ☐ Fit the model using `aov()` and report F , df , and P .
- ☐ Run post-hoc comparisons using `emmeans()`.
- ☐ Write a conclusion that addresses the null hypothesis and explains the results in a biological context.

Conclusion

Closing thoughts

In this lab we worked through the full one-way ANOVA workflow: exploring grouped data, fitting a model with `aov()`, checking assumptions, running post-hoc comparisons with `emmeans`, and reporting results. We also saw that a two-sample t -test and a one-way ANOVA with two groups are mathematically equivalent ($F = t^2$, same P-value and residual degrees of freedom).

Next week we will look at model residuals in more detail and learn formal diagnostic checks for the ANOVA assumptions.

Attribution

This lab was developed using resources that are available under a [Creative Commons Attribution 4.0 International license](#), made available on the [SOLES Open Educational Resources repository](#).