

# ENVX2001 Lab 07 - Regression model development

ENVX2001 Applied Statistical Methods  
Semester 1, 2026

## Learning outcomes

In this lab, you will work towards achieving learning outcomes L03, and L05.

### Lab objectives

In this lab, we will:

- Identify best predictors for model - Exercise 1
- Fit model and check assumptions - Exercise 1
- Interpret model output - Exercise 1



Tip

Please work on this exercise by creating your own R Markdown file.

## Preparation

- Install or update the `performance` package

CODE


```
#install.packages("performance")
library(performance)
```

This package is really good for checking your models. For this lab, we will focus on the `check_model()` function, which gives us nice pretty diagnostic plots for models:

### `plot()`

CODE

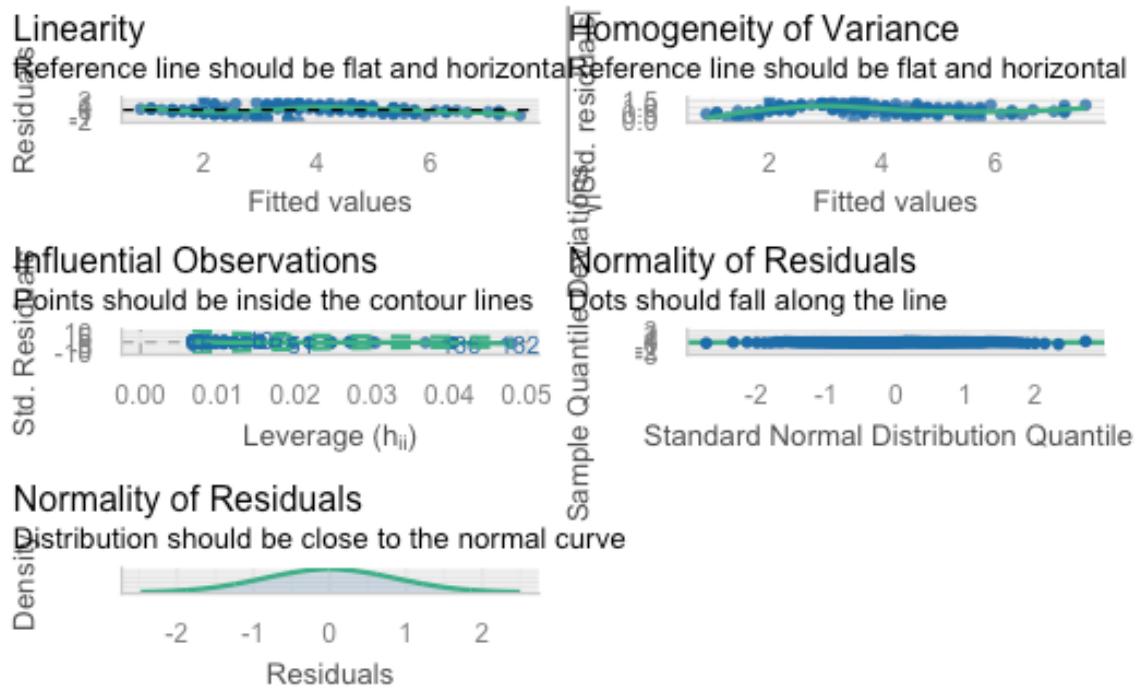
```
CODE  
par(mfrow=c(2,2))  
plot(iris_lm)
```



```
CODE  
par(mfrow=c(1,1))
```

## check\_model() from performance

```
CODE  
library(performance)  
check_model(iris_lm)
```



## Exercise 1: Modelling bird abundance

We will now use the transformed data in `loyn` for this exercise. If you have not already figured out how to perform the transformation, or if something is wrong, you may use the `loyn` tab in the `mlr.xlsx` MS Excel document. Alternatively, the code to convert the data is below.

### 💡 Tip

This is the same data we used in the walkthrough exercise

```
CODE
# Load library if needed
library(readxl)
# reset the data import just in case it has been modified
loyn ← read_xlsx("data/mlr.xlsx", "Loyn")
# make transformations

loyn ← loyn %>%
  mutate(
    L10AREA = log10(AREA),
```

```

L10DIST = log10(DIST),
L10LDIST = log10(LDIST)
)

# check
glimpse(loyn)

```

#### OUTPUT

```

Rows: 56
Columns: 10
$ ABUND    <dbl> 5.3, 2.0, 1.5, 17.1, 13.8, 14.1, 3.8, 2.2, 3.3, 3.0, 27.6, 1...
$ AREA     <dbl> 0.1, 0.5, 0.5, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 2.0, 2.0, 2...
$ YR.ISOL  <dbl> 1968, 1920, 1900, 1966, 1918, 1965, 1955, 1920, 1965, 1900, 1...
$ DIST     <dbl> 39, 234, 104, 66, 246, 234, 467, 284, 156, 311, 66, 93, 39, 4...
$ LDIST    <dbl> 39, 234, 311, 66, 246, 285, 467, 1829, 156, 571, 332, 93, 39, ...
$ GRAZE   <dbl> 2, 5, 5, 3, 5, 3, 5, 4, 5, 3, 5, 2, 1, 5, 5, 3, 3, 2, 2...
$ ALT      <dbl> 160, 60, 140, 160, 140, 130, 90, 60, 130, 210, 160, 210, ...
$ L10AREA  <dbl> -1.0000000, -0.3010300, -0.3010300, 0.0000000, 0.0000000, 0.0...
$ L10DIST  <dbl> 1.591065, 2.369216, 2.017033, 1.819544, 2.390935, 2.369216, 2...
$ L10LDIST <dbl> 1.591065, 2.369216, 2.492760, 1.819544, 2.390935, 2.454845, 2...

```

## Best single predictor?

### Question 1

Obtain the correlation between ABUND and all of the predictor variables using `cor()`. Based on these, what would you expect to be the best single predictor of ABUND?

**CODE**

```
cor(loyn)
```

## Assumptions and interpretation

### Question 2

Use multiple linear regression to see whether ABUND can be predicted from L10AREA and GRAZE. Are the assumptions met? Is there a significant relationship? *Note: we are using these 2 predictors as they have the largest absolute correlations. Use `lm()` and specify the model as `ABUND ~ L10AREA + GRAZE`.*

**CODE**

```

lm.mod1 ← lm(ABUND ~ GRAZE + L10AREA, data = loyn)

par(mfrow = c(2, 2))
plot(lm.mod1)
par(mfrow = c(1, 1))

summary(lm.mod1)

```

## Question 3

How good is the model based on the (i)  $r^2$  (ii) adjusted  $r^2$ ? Use `summary()`.

CODE

```
summary(lm.mod1)$r.squared  
summary(lm.mod1)$adj.r.squared
```

## Question 4

Which variable(s) has the most significant effect(s)? (*Refer specifically to the t probabilities in the table of predictors and their estimated parameters or coefficients in the output of `summary()`.*). Interpret the p-values in terms of dropping predictor variables.

## Question 5

Repeat the multiple regression, but this time include YRS.ISOL as a predictor variable (it has the 3rd largest absolute correlation). This will allow you to assess the effect of YRS.ISOL with the other variables taken into account.

## Question 6

Check assumptions, do the residuals look ok? If you are happy with the assumptions, you can proceed to interpret the model output.

## Question 7

Compare the  $r^2$  and adjusted  $r^2$  values with those you calculated for the 2 predictor model, Which is the better model? Why?

CODE

```
summary(lm.mod2)
```

# Exercise 2: California streamflow

The following dataset contains 43 years of annual precipitation measurements (in mm) taken at (originally) 6 sites in the Owens Valley in California. I have reduced this to three variables labelled `lake_sabrina` (Lake Sabrina), `pine_creek` (Big Pine Creek), `rock_creek` (Rock Creek), and the dependent variable stream runoff volume (measured in ML/year) at a site near Bishop, California (labelled `runoff_volume`).

Note the variables have already been log-transformed to increase normality of the residuals in the regressions.

Start with a full model and manually remove the variables one at a time, checking every time whether removal of a variable actually improves the model.

```
CODE
# read in the data
stream_data <- read_xlsx("data/california_streamflow.xlsx", "streamflow")
names(stream_data)
```

```
OUTPUT
[1] "lake_sabrina"   "pine_creek"      "rock_creek"     "runoff_volume"
```

```
CODE
s.mod_full <- lm(runoff_volume~lake_sabrina + pine_creek + rock_creek, data=stream_data)
s.mod_full <- lm(runoff_volume~., data=stream_data) ## you can also use the . to indicate use all
variables
summary(s.mod_full)
```

```
OUTPUT
Call:
lm(formula = runoff_volume ~ ., data = stream_data)

Residuals:
    Min      1Q  Median      3Q      Max 
-0.09885 -0.03331  0.01025  0.03359  0.09495 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 3.25716   0.12360  26.352 < 2e-16 ***
lake_sabrina 0.05631   0.03756  1.499  0.14185    
pine_creek   0.21085   0.06756  3.121  0.00339 **  
rock_creek   0.43838   0.08798  4.983 1.32e-05 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.04861 on 39 degrees of freedom
Multiple R-squared:  0.8817,    Adjusted R-squared:  0.8726 
F-statistic: 96.88 on 3 and 39 DF,  p-value: < 2.2e-16
```

## Partial F-Tests

The above analysis tells us that both `pine_creek` & `rock_creek` are significant, according to the t-test, in the model and `lake_sabrina` is not? This involves performing Partial F-Tests as discussed in the lecture.

This can be done in **R** by using `anova()` on two model objects. To be able to compare the models and run the `anova`, you need to make objects of all the possible model combinations you want to compare.

```
CODE
```

```
s.mod_reduced ← lm(runoff_volume ~ rock_creek + pine_creek, data=stream_data)
anova(s.mod_reduced, s.mod_full)
```

The last row gives the results of the partial F-test.

## Question 1

Should we remove lake\_sabrina from the model?

## Question 2

Is the p-value for the f-test the same as for the t-test?

## Question 3

Write out the hypotheses you are testing.

Perform a Partial F-Test to work out if the removal of lake\_sabrina and pine\_creek improves upon the full model.

```
CODE
s.mod_reduced2 ← lm(runoff_volume ~ lake_sabrina + pine_creek, data=stream_data)
anova(s.mod_reduced2, s.mod_full)
```

OUTPUT

```
Analysis of Variance Table

Model 1: runoff_volume ~ lake_sabrina + pine_creek
Model 2: runoff_volume ~ lake_sabrina + pine_creek + rock_creek
  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     40 0.150845
2     39 0.092166  1  0.05868 24.83 1.321e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Question 4

Which variable should be added to the model containing rock\_creek?

## Question 5

Could things be even simpler? Perform a partial F-Test to see if a model containing rock\_creek alone could be suitable.

```
CODE
s.mod_reduced3 ← lm(runoff_volume ~ rock_creek, data=stream_data)
anova(s.mod_reduced3, s.mod_full)
```

## Question 6

What is your optimal model?

# Review

- Simple linear regressions model the relationship between two variables
  - ▶ We can also make linear models with more than one predictor
- We can use histograms and correlation matrices to do some preliminary exploration of the data
- If any variables are skewed, we can transform them
- Looking at a correlation matrix to identify the best predictors (for both simple and multiple linear regression)
- Fit model using `lm()` function
- Check assumptions:
  - ▶ Collinearity (multiple linear regression only)
  - ▶ Linearity
  - ▶ Independence
  - ▶ Normality
  - ▶ Equal variance
- Use `summary()` to look at model output and interpret it
  - ▶ F-test : overall model significance
  - ▶ Coefficients table : individual predictors' significance
  - ▶  $R^2$  : How much variation in the data is explained by the model?

That's it for today! Great work fitting simple and multiple linear regression! Next week we jump into stepwise selection!

## Attribution