

Lab 06

ENVX2001 Applied Statistical Methods

Contents

Welcome!	2
Learning outcomes	2
Specific goals	2
Preparation	2
Downloads	2
1. Factorial ANOVA as CRD (~30 min)	3
Analysis	4
Exercise 1	4
Solution	4
(i) The statistical model for the factorial ANOVA is:	4
(ii) Test the assumptions of the factorial ANOVA and transform the data (if needed).	5
(iii) Use an ANOVA to analyse the CRD and report on the results.	7
(iv) Explore any significant results using post-hoc tests from emmeans package and its emmeans() function. Include any plots that you think are useful to visualise the results.	8
2. Factorial ANOVA with blocking (~30 min)	10
Analysis	11
Exercise 2	11
Solution	11
(i) The statistical model for the factorial ANOVA with blocking is:	11
(ii) Test the assumptions of the ANOVA model and transform the data (if needed).	12
(iii) Use an ANOVA to analyse the RCBD and report on the results.	13
(iv) Calculate how much variation the blocking term captured.	13
Conclusion	14
Closing thoughts	14
Attribution	14

Welcome!

This week we will analyse a factorial experiment using ANOVA: first as a completely randomised design (CRD), then as a randomised complete block design (RCBD). The experiment investigates aggressive interactions in Tasmanian devils.

You will come across **Exercises** in this lab. Solutions will be posted on Friday evening.

Learning outcomes

At the end of this practical students should be able to:

- use R to analyse experiments with a factorial treatment structure where the experimental design is a CRD or RCBD;

Specific goals

By the end of this lab, you should be able to:

- ☐ Fit and interpret a factorial ANOVA for a CRD
- ☐ Check ANOVA assumptions and apply transformations when needed
- ☐ Fit a factorial ANOVA with a blocking term for an RCBD
- ☐ Perform post-hoc tests using emmeans
- ☐ Calculate the percentage of variation explained by blocking

Preparation

This lab uses emmeans. Install it if you are missing it by running the following **in the console**:

```
CODE  
install.packages("emmeans")
```

Downloads

File	Used in	Download
tasmanian_devil_aggression.csv	Sections 1 and 2	Download

1. Factorial ANOVA as CRD (~30 min)

You are interested in investigating if aggressive interactions between Tasmanian devil individuals are influenced by eating in a group or alone, and if they are male or female. You conduct an experiment where you observe the number of aggressive interactions between individuals in a 10 minute trial. You have 12 nights to conduct the experiment and each night you observe one trial for each of the 4 treatment combinations.

Analyse the data using a factorial ANOVA as if it was collected by a CRD.

The data is in the **tasmanian_devil_aggression.csv** file.

Factors:

- feeding_context: Alone vs Group
- sex: Female vs Male

Response variables:

- aggression_count (integer count per trial)
- log_aggression = $\log(\text{aggression_count} + 1)$

Blocking factor (we will ignore this for the CRD analysis but it will be used in the RCBD analysis):

- block: Night_01 ... Night_12 (each night includes all 4 treatment combinations)
- Total N: $12 \text{ blocks} \times 2 \text{ sexes} \times 2 \text{ contexts} = 48 \text{ trials}$ (1 obs per cell per block)



Figure 1: Tasmanian devil (*Sarcophilus harrisii*) is the largest carnivorous marsupial in the world.

Image credit: Wayne McLean.

Analysis

Exercise 1

(i) Write out the statistical model for the factorial ANOVA for this experiment. Define all the terms in the model. You can write out in words or in mathematical notation.

```
CODE
# Answer here
```

(ii) Test the assumptions of the factorial ANOVA and transform the data (if needed). Hint first run the model using the `aov()` function and then use the `plot()` function to check the assumptions.

```
CODE
#
```

(iii) Use an ANOVA to analyse the CRD and report on the results.

```
CODE
#
```

(iv) Explore any significant results using post-hoc tests from `emmeans` package and its `emmeans()` function. Include any plots that you think are useful to visualise the results.

```
CODE
#
```

Solution

(i) **The statistical model for the factorial ANOVA is:**

$$Y_{ijk} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \epsilon_{ijk}$$

Where:

- Y_{ijk} is the response variable (aggression count) for the k th observation in the i th level of feeding context and j th level of sex
- μ is the overall mean aggression count
- α_i is the effect of the i th level of feeding context (Alone vs Group)
- β_j is the effect of the j th level of sex (Female vs Male)
- $(\alpha\beta)_{ij}$ is the interaction effect between feeding context and sex

- ϵ_{ijk} is the random error term for the k th observation in the i th level of feeding context and j th level of sex and is assumed to be normally distributed with mean 0 and constant variance σ^2 .

OR:

Aggression count = overall mean + feeding context + sex + feeding context $\times$ sex + random error

(ii) Test the assumptions of the factorial ANOVA and transform the data (if needed).

Read in the data and convert to factor if required.

CODE

```
devil <- read.csv("data/tasmanian_devil_aggression.csv")
str(devil)
```

OUTPUT

```
'data.frame': 48 obs. of 6 variables:
 $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ block   : chr  "Night_01" "Night_01" "Night_01" "Night_01" ...
 $ sex     : chr  "Female" "Female" "Male" "Male" ...
 $ feeding_context : chr  "Alone" "Group" "Alone" "Group" ...
 $ aggression_count: int  3 5 5 8 1 4 3 3 5 3 ...
 $ log_aggression : num  1.386 1.792 1.792 2.197 0.693 ...
```

CODE

```
devil$feeding_context <- as.factor(devil$feeding_context)
devil$sex <- as.factor(devil$sex)
str(devil)
```

OUTPUT


```
'data.frame': 48 obs. of 6 variables:
 $ id      : int  1 2 3 4 5 6 7 8 9 10 ...
 $ block   : chr  "Night_01" "Night_01" "Night_01" "Night_01" ...
 $ sex     : Factor w/ 2 levels "Female","Male": 1 1 2 2 1 1 2 2 1 1 ...
 $ feeding_context : Factor w/ 2 levels "Alone","Group": 1 2 1 2 1 2 1 2 1 2 ...
 $ aggression_count: int  3 5 5 8 1 4 3 3 5 3 ...
 $ log_aggression : num  1.386 1.792 1.792 2.197 0.693 ...
```

Run the model and check the assumptions with the `plot()` function.

CODE

```
devil.aov <- aov(aggression_count ~ feeding_context*sex, data=devil)

par(mfrow=c(1,2))
plot(devil.aov, which=1)
plot(devil.aov, which=2)
```



CODE


```
par(mfrow=c(1,1))
```

The residual diagnostics show that the data is not normally distributed, has unequal variance (fanning). A log transformation is applied to the response variable to try and meet the assumptions.

CODE

```
devil.log.aov <- aov(log(aggression_count+1) ~ feeding_context*sex, data=devil)

par(mfrow=c(1,2))
plot(devil.log.aov, which=1)
plot(devil.log.aov, which=2)
```



CODE

```
par(mfrow=c(1,1))
```

The residual diagnostics show that the data is not normally distributed, but now has equal variance. Given ANOVA is robust against non-normality, we will proceed with the analysis using the log transformed response variable.

(iii) Use an ANOVA to analyse the CRD and report on the results.

CODE

```
summary(devil.log.aov)
```

OUTPUT

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
feeding_context	1	4.937	4.937	16.637	0.000187	***
sex	1	2.752	2.752	9.274	0.003917	**
feeding_context:sex	1	0.082	0.082	0.276	0.601716	
Residuals	44	13.058	0.297			

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The ANOVA table shows that there is not a significant interaction between `feeding_context` and `sex` with a P-value of 0.60. This means we can now move to the main effects and see that both `feeding_context` and `sex` have a significant influence on `log_aggression`.

(iv) Explore any significant results using post-hoc tests from emmeans package and its emmeans() function. Include any plots that you think are useful to visualise the results.

Main effect of feeding_context is significant so we can use emmeans() to conduct post-hoc tests to determine which pairs are significantly different. The results show that the log_aggression is significantly higher when feeding in a group compared to alone.

CODE

```
library(emmeans)
```

OUTPUT

```
Welcome to emmeans.  
Caution: You lose important information if you filter this package's results.  
See '? untidy'
```

CODE

```
emmeans(devil.log.aov, pairwise ~ feeding_context)
```

OUTPUT

```
NOTE: Results may be misleading due to involvement in interactions
```

OUTPUT

```
$emmeans  
feeding_context emmean    SE df lower.CL upper.CL  
Alone           1.12 0.111 44    0.894    1.34  
Group           1.76 0.111 44    1.535    1.98
```

```
Results are averaged over the levels of: sex  
Results are given on the log(mu + 1) (not the response) scale.  
Confidence level used: 0.95
```

```
$contrasts  
contrast      estimate    SE df t.ratio p.value  
Alone - Group -0.641 0.157 44  -4.079 0.0002
```

```
Results are averaged over the levels of: sex  
Note: contrasts are still on the log(mu + 1) scale. Consider using  
    regrid() if you want contrasts of back-transformed estimates.
```


Plotting the marginal means for feeding_context shows that the log_aggression is higher when feeding in a group compared to alone.

CODE

```
plot(emmeans(devil.log.aov, pairwise ~ feeding_context), comparisons = TRUE)
```

OUTPUT

```
NOTE: Results may be misleading due to involvement in interactions
```

The main effect of sex is significant so we can use `emmeans()` to conduct post-hoc tests to determine which pairs are significantly different. The results show that the `log_aggression` is significantly

CODE

```
emmeans(devil.log.aov, pairwise ~ sex)
```

OUTPUT

NOTE: Results may be misleading due to involvement in interactions

OUTPUT

```
$emmeans
  sex    emmean    SE df lower.CL upper.CL
Female 1.20 0.111 44    0.975    1.42
Male   1.68 0.111 44    1.454    1.90
```

Results are averaged over the levels of: feeding_context
Results are given on the log($\mu + 1$) (not the response) scale.
Confidence level used: 0.95

```
$contrasts
 contrast      estimate    SE df t.ratio p.value
Female - Male -0.479 0.157 44   -3.045 0.0039
```

Results are averaged over the levels of: feeding_context
Note: contrasts are still on the log($\mu + 1$) scale. Consider using
`regrid()` if you want contrasts of back-transformed estimates.

Males have a significantly higher `log_aggression` compared to females.


Plotting:

CODE

```
plot(emmeans(devil.log.aov, pairwise ~ sex), comparisons = TRUE)
```

OUTPUT

NOTE: Results may be misleading due to involvement in interactions



Before we move on, now is a good time to take a 5-minute break.

2. Factorial ANOVA with blocking (~30 min)

Now we will analyse the same experiment as above but this time we will include the blocking factor of `block` (night) in the analysis. This means we will be analysing the data as if it was collected by a RCBD.

Analysis

Exercise 2

(i) Write out the statistical model for the factorial ANOVA with blocking for this experiment. Define all the terms in the model. You can write out in words or in mathematical notation.

```
CODE
# Answer here
```

(ii) Test the assumptions of the ANOVA model and transform the data (if needed). Hint first run the model using the `aov()` function and then use the `plot()` function to check the assumptions.

```
CODE
#
```

(iii) Use an ANOVA to analyse the RCBD and report on the results.

```
CODE
#
```

(iv) Calculate how much variation the blocking term captured.

```
CODE
#
```

Solution

(i) **The statistical model for the factorial ANOVA with blocking is:**

$$Y_{ijkl} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \gamma_k + \epsilon_{ijkl}$$

Where:

- Y_{ijkl} is the response variable (aggression count) for the l th observation in the i th level of feeding context, j th level of sex and k th block (night)
- μ is the overall mean aggression count
- α_i is the effect of the i th level of feeding context (Alone vs Group)
- β_j is the effect of the j th level of sex (Male vs Female)
- $(\alpha\beta)_{ij}$ is the interaction effect between feeding context and sex
- γ_k is the effect of the k th block (night)
- ϵ_{ijkl} is the random error term for the l th observation in the i th level of feeding context, j th level of sex and k th block (night) and is assumed to be normally distributed with mean 0 and constant variance σ^2 .


OR:

$Aggression\ count = overall\ mean + block + feeding\ context + sex + feeding\ context \times sex + random\ error$

(ii) Test the assumptions of the ANOVA model and transform the data (if needed).

Run the model and check the assumptions with the `plot()` function.


```
CODE
devil.block.aov <- aov(aggression_count ~ block + feeding_context * sex, data=devil)
par(mfrow=c(1,2))
plot(devil.block.aov, which=1)
plot(devil.block.aov, which=2)
```



```
CODE
par(mfrow=c(1,1))
```

The residual diagnostics show that the data is not normally distributed, but has equal variance. A log transformation is applied to the response variable to try and meet the assumption of normality.

```
CODE
devil.block.log.aov <- aov(log(aggression_count+1) ~ block + feeding_context * sex,
data=devil)
par(mfrow=c(1,2))
plot(devil.block.log.aov, which=1)
plot(devil.block.log.aov, which=2)
```



```
CODE
par(mfrow=c(1,1))
```

The residual diagnostics show that the data is close to being normally distributed. Given ANOVA is robust against non-normality, we will proceed with the analysis using the log transformed response variable.

(iii) Use an ANOVA to analyse the RCBD and report on the results.

```
CODE
summary(devil.block.log.aov)
```

```
OUTPUT
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
block	11	4.115	0.374	1.380	0.227950
feeding_context	1	4.937	4.937	18.219	0.000156 ***
sex	1	2.752	2.752	10.156	0.003140 **
feeding_context:sex	1	0.082	0.082	0.303	0.585914
Residuals	33	8.943	0.271		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The ANOVA table shows that there is not a significant interaction between feeding_context and sex with a P-value of 0.59. However, both main effects are significant.

(iv) Calculate how much variation the blocking term captured.

The sum of squares for the block term is 4.115 Total sum of squares is:

CODE

```
Total.SS ← 4.115+4.937+2.752+0.082+8.943
```

```
Total.SS
```

OUTPUT

```
[1] 20.829
```

The percentage of variation captured by the block term is:

CODE

```
(4.115/Total.SS)*100
```

OUTPUT

```
[1] 19.75611
```

The block term captured 19.76% of the variation in the data, so it was useful using a RCBD for this experiment.

Conclusion

Closing thoughts

We analysed a factorial experiment this week using Tasmanian devil aggression data: first as a CRD to test for main effects and interactions, then as an RCBD to account for night-to-night variation. These factorial designs and analysis techniques extend the one-way ANOVA methods from previous weeks.

Attribution

This lab was developed using resources that are available under a [Creative Commons Attribution 4.0 International license](#), made available on the [SOLES Open Educational Resources repository](#).