

Tutorial 08

ENVX2001 – Applied Statistical Methods

Semester 1

from tut07 old - need to use this to update for week 8!!:

Bird abundance

Fragmentation of forest habitat has an impact of wildlife abundance. This study looked at the relationship between bird abundance (bird ha⁻¹) and the characteristics of forest patches at 56 locations in SE Victoria.

The predictor variables are:

- ALT Altitude (m)
- YR.ISOL Year when the patch was isolated (years)
- GRAZE Grazing (coded 1-5 which is light to heavy)
- AREA Patch area (ha)
- DIST Distance to nearest patch (km)
- LDIST Distance to largest patch (km)

Import the data from the “Loyn” tab in the MS Excel file.

CODE

```
library(readxl)
loyn <- read_xlsx("data/mlr.xlsx", "Loyn")
```

Often, the first step in model development is to examine the data. This is a good way to get a feel for the data and to identify any issues that may need to be addressed. In this case, we will examine the data using histograms and a correlation matrix.


Histograms

There are a breadth of ways to create histograms in R. In each tab below you will find some different ways to create the same plot outputs.

hist()


This is a straightforward way to create multiple histograms with `hist()`. The `par()` function is used to arrange the plots on the page. The `mfrow` argument specifies the number of rows and columns of plots.

```
CODE
# par(mfrow=c(3,3))
hist(loyn$ABUND)
```



```
CODE
hist(loyn$ALT)
```


Histogram of loyn\$ALT



CODE


```
hist(loyn$YR.ISOL)
```

Histogram of loyn\$YR.ISOL



CODE


```
hist(loyn$GRAZE)
```



CODE

```
hist(loyn$AREA)
```


Histogram of loyn\$AREA



CODE


```
hist(loyn$DIST)
```

Histogram of loyn\$DIST



CODE

```
hist(loyn$LDIST)
```



CODE

```
# par(mfrow=c(1,1))
```

hist.data.frame() from Hmisc

The Hmisc package provides a function `hist.data.frame()` that can be used to create multiple histograms, which can be called by simply using `hist()`. You may need to tweak the `nclass` argument to get the desired number of bins, as the default may not look appropriate.

CODE

```
# install.packages("Hmisc")  
library(Hmisc)  
hist(loyn, nclass = 50)
```


ggplot()

A more modern approach is to use `ggplot()` with `facet_wrap()` to arrange multiple plots on a single page. To do this, the `pivot_longer()` function from the `tidyr` package is used to reshape the data into a tidy format.

CODE

```
# tidy the data
loyn_tidy <- pivot_longer(loyn, cols = everything())


# plot
ggplot(loyn_tidy, aes(x = value)) +
  geom_histogram() +
  facet_wrap(~name, scales = "free") +
  theme_bw()
```



ggplot() with dplyr

Here we use the pipe operator `%>%` from `dplyr` to chain together a series of commands. The pipe operator takes the output of the command on the left and passes it to the command on the right (or below) the pipe. This means that we can create a series of commands that are executed in order.

```
CODE
loyn %>%
  pivot_longer(cols = everything()) %>%
  ggplot(aes(x = value)) +
  geom_histogram() +
  facet_wrap(~name, scales = "free") +
  theme_bw()
```



Question 1

Comment on the histograms in terms of leverage. *Hint: what is the relationship between leverage and skewness?*

Correlation matrix

Calculate the correlation matrix using `cor(Loyn)`.

CODE

```
cor(Loyn)
```

OUTPUT

	ABUND	AREA	YR.ISOL	DIST	LDIST
ABUND	1.00000000	0.255970206	0.503357741	0.2361125	0.08715258
AREA	0.25597021	1.000000000	-0.001494192	0.1083429	0.03458035
YR.ISOL	0.50335774	-0.001494192	1.000000000	0.1132175	-0.08331686
DIST	0.23611248	0.108342870	0.113217524	1.0000000	0.31717234
LDIST	0.08715258	0.034580346	-0.083316857	0.3171723	1.00000000
GRAZE	-0.68251138	-0.310402417	-0.635567104	-0.2558418	-0.02800944
ALT	0.38583617	0.387753885	0.232715406	-0.1101125	-0.30602220
GRAZE					
ALT					
ABUND	-0.68251138	0.3858362			
AREA	-0.31040242	0.3877539			
YR.ISOL	-0.63556710	0.2327154			
DIST	-0.25584182	-0.1101125			
LDIST	-0.02800944	-0.3060222			
GRAZE	1.00000000	-0.4071671			
ALT	-0.40716705	1.0000000			

Question 2

Which independent variables are useful for predicting the dependent variable abundance? Is there evidence for multi-collinearity?


Plotting correlation

Examine correlations visually using `pairs()` or `corrplot()` from the `corrplot` package.

Scatterplot matrix

CODE


```
pairs(loyn)
```



Correlation matrix

CODE

```
library(corrplot)
corrplot(cor(loyn))
```



Question 3

Are there any trends visible from the plots?

Tip

We can also bring in variance inflation factors (VIF) to help us identify multi-collinearity, but that is done only after we have selected a model.

Transformations

The AREA predictor has a small number of observations with very large values. Apply a $\log_{10}$ transformation and label the new variable `Loyn$L10AREA`.

CODE

```
Loyn$L10AREA <- log10(Loyn$AREA)
```

Question 4

Why are we transforming AREA?

Question 5

Re-run `pairs(Loyn)` and create a histogram using the transformed value of AREA, how do the plots look?

CODE

```
hist(loyn$L10AREA)  
pairs(loyn)
```

Question 6

In preparation for modelling, transform the remaining skewed variables, DIST and LDIST the same way you did for AREA and examine the histogram and pairs plots using these new variables.

Make sure you end up with two new variables labelled `loyn$L10DIST` and `loyn$L10LDIST`.

i Checkpoint

You should now be able to ...[PUT THIS AFTER EACH OBJECTIVE]

Wrap-up

In this tutorial, we explored factorial designs and their analysis using ANOVA. We learned about main effects, interactions, and blocking effects, and how to interpret ANOVA results in the context of factorial experiments. We also visualised these effects using `ggplot2` and `emmeans`, enhancing our understanding of the relationships between factors and response variables in experimental designs.

Example exam questions

1. In a 2-way factorial design with factors A (2 levels) and B (3 levels), describe how you would identify and interpret main effects and interactions using ANOVA. Provide an example of how to visualise these effects.
2. Explain the difference between a treatment design and an experimental design. Provide examples of each and discuss how they can be combined in a factorial design with blocking.
3. If you conducted a factorial experiment with blocking and found a significant interaction between two factors, how would you interpret this result? Would you still consider the main effects of each factor? Justify your answer with reference to ANOVA principles.

Exercise 1: Backward elimination in R

Data: *Dippers* spreadsheet

Dippers are thrush-sized birds living mainly in the upper reaches of rivers, which feed on benthic invertebrates by probing the river beds with their beaks. The dataset in this exercise contains data from a biological survey which examined the nature of the variables thought to influence the breeding of British dippers.

Twenty-two sites were included in the survey. Some of the variables have been transformed.

The variables measured were:

- Altitude site altitude
- Hardness water hardness
- RiverSlope river-bed slope
- LogCadd the numbers of caddis fly larvae, transformed
- LogStone the numbers of stonefly larvae, transformed
- LogMay the numbers of mayfly larvae, transformed
- LogOther the numbers of all other invertebrates collected, transformed
- Br_Dens the number of breeding pairs of dippers per 10 km of river

In the analyses, the four invertebrate variables were transformed using a $\log(x + 1)$ transformation.

CODE

```
library(readxl)
Dippers <- read_xlsx("data/mlr.xlsx" , sheet = "Dippers")
glimpse(Dippers)
```

OUTPUT

```
Rows: 22
Columns: 8
$ Altitude <dbl> 259, 198, 251, 184, 145, 145, 198, 160, 251, 159, 160, 145,...
$ Hardness <dbl> 12.20, 22.00, 26.30, 22.50, 29.50, 39.90, 42.80, 59.60, 69.0...
$ RiverSlope <dbl> 10.90, 14.70, 6.90, 4.60, 1.91, 5.00, 6.20, 14.30, 4.60, 3.0...
$ Br_Dens <dbl> 3.60, 4.30, 3.80, 3.40, 3.80, 4.50, 4.30, 5.00, 4.50, 3.40,...
$ LogCadd <dbl> 2.303, 2.890, 3.784, 4.419, 3.219, 3.932, 3.664, 4.431, 3.7...
$ LogStone <dbl> 5.242, 4.344, 5.231, 5.242, 3.829, 4.898, 4.357, 6.337, 5.4...
$ LogMay <dbl> 0.000, 3.401, 5.826, 5.749, 5.509, 5.749, 5.371, 0.000, 4.8...
$ LogOther <dbl> 1.386, 1.609, 1.386, 1.386, 1.099, 3.045, 1.386, 2.944, 2.5...
```

You may explore the data on your own (hint: look at last week's exercise on histogram and scatterplot matrices).

When ready, perform a backward elimination starting from the full model:

CODE

```
FullMod <- lm(Br_Dens ~ ., data=Dippers)
RedMod <- step(FullMod, direction = "backward")
summary(FullMod)
summary(RedMod)
AIC(FullMod, RedMod)
```

Question 1

Which model is chosen? Why?